

UNIVERSITY OF PENNSYLVANIA  
ESE 650: LEARNING IN ROBOTICS  
SPRING 2023

[03/13] HOMEWORK 3

DUE: 04/10 MON 11.59 PM ET

---

**Changelog:** This space will be used to note down updates/errata to the homework problems.

- **Page 4: Line 21: Third bullet is removed. In short the robot can enter obstacles (which are all grey cells) but it has to stay there indefinitely and incur the penalty.**

---

Read the following instructions carefully before beginning to work on the homework.

- You will submit solutions typeset in  $\text{\LaTeX}$  on Gradescope (strongly encouraged). You can use `hw_template.tex` on Canvas in the “Homeworks” folder to do so. If your handwriting is *unambiguously legible*, you can submit PDF scans/tablet-created PDFs.
- Please start a new problem on a fresh page and mark all the pages corresponding to each problem. Failure to do so may result in your work not graded completely.
- Clearly indicate the name and Penn email ID of all your collaborators on your submitted solutions.
- **For each problem in the homework, you should mention the total amount of time you spent on it. This helps us gauge the perceived difficulty of the problems.**
- You can be informal while typesetting the solutions, e.g., if you want to draw a picture feel free to draw it on paper clearly, click a picture and include it in your solution. Do not spend undue time on typesetting solutions.
- You will see an entry of the form “HW 3 PDF” where you will upload the PDF of your solutions. You will also see entries like “HW 3 Problem 1 Code” where you will upload your solution for the respective problems. **For each programming problem, you should create a fresh Python file.** This file should contain all the code to reproduce the results of the problem and you will upload the .py file to Gradescope. If we have installed Autograder

for a particular problem, you will use the Autograder. Name your file to be “pennkey\_hw3\_problem1.py”, e.g., I will name my code for Problem 1 as “pratikac\_hw3\_problem1.py”.

- **You should include all the relevant plots in the PDF, without doing so you will not get full credit. There is no auto-grader for this homework so this is particularly important.** You can, for instance, export your Jupyter notebook as a PDF (you can also use text cells to write your solutions) and export the same notebook as a Python file to upload your code.
- **Your PDF solutions should be completely self-contained. We will run the Python file to check if your solution reproduces the results in the PDF.**

Credit. The points for the problems add up to 110. You only need to solve for 100 points to get full credit, i.e., your final score will be  $\min(\text{your total points}, 100)$ .

---


1 **Problem 1 (Policy Iteration, 20 points).** Consider the following Markov Decision  
2 Process. The state-space is a  $10 \times 10$  grid, cells that are obstacles are marked in gray.  
3 The initial state of the robot is in blue and our desired terminal state is in green. The  
4 robot gets a *reward* of 10 if it reaches the desired terminal state with a discount  
5 factor of 0.9. At each non-obstacle cell, the robot can attempt to move to any of the  
6 immediate neighboring cells using one of the four controls (North, East, West and  
7 South). The robot cannot diagonally. The move succeeds with probability 0.7 and  
8 with remainder probability 0.3 the robot can end up at some other cell as follows:

$$P(\text{moves north} \mid \text{control is north}) = 0.7,$$

$$P(\text{moves west} \mid \text{control is north}) = 0.1,$$

$$P(\text{moves east} \mid \text{control is north}) = 0.1,$$

$$P(\text{does not move} \mid \text{control is north}) = 0.1.$$



9

10 Similarly, if the robot desired to go east, it may end up in the cells to its north, south,  
11 or stay put at the original cell with total probability 0.3 and actually move to the  
12 cell east with probability 0.7. The cost pays a cost of 1 (i.e., reward is -1) for each  
13 control input it takes, regardless of the outcome. If the robot ends up at a state  
14 marked as an obstacle (all grey cells are obstacles, i.e., cell marked (0,0), (0,1), (3,2)  
15 etc. are obstacles), it gets a reward of -10 for each time-step that it remains inside  
16 the obstacle cell. The robot is allowed to stay in the goal state indefinitely (i.e., take  
17 a special action to “not move”) and this action gets no reward/cost.

18 We would like to implement policy iteration to find the best trajectory for the  
19 robot to go from the blue cell to the green cell.

20 (a) **(0 points)** Carefully code up the above environment to run policy iteration.  
21 You will need to think about how to code up the probability transition matrix  
22  $\mathbb{R}^{100 \times 100} \ni T_{x,x'}(u) = P(x' \mid x, u)$ , the run-time cost  $q(x, u)$ , and the  
23 terminal cost  $q_f(x)$ . Policy iteration is easy to implement if you represent  
24 all the above quantities as matrices and vectors. Plot the environment to  
25 check if it confirms to the above picture.

- 1 (b) **(10 points)** Initialize policy iteration with a feedback control  $u^{(0)}(x)$  where  
 2 the robot always goes east, this results in a policy  $\pi^{(0)} = (u^{(0)}(\cdot), u^{(0)}(\cdot), \dots)$ .  
 3 Write the code for policy evaluation to obtain the cost-to-go from every cell  
 4 in the above picture for this initial policy. Plot the value function  $J^{\pi^{(0)}}(x)$   
 5 as a heatmap in the above picture.
- 6 (c) **(10 points)** Execute the policy iteration algorithm, you will iteratively  
 7 perform policy evaluation and policy improvement steps. For the first 4 iter-  
 8 ations, plot the feedback control  $u^{(k)}(x)$  (using arrows as shown in the lecture  
 9 notes ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.arrow.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.arrow.html)),  
 10 you can also write the control input in the cell). You should color the cell  
 11 using the value function  $J^{\pi^{(k)}}(x)$ .


12 We have left the transition probabilities and the reward structure a bit vague to  
 13 force you to think carefully of the nuances of this problem. But some clarification  
 14 could be useful.

- 15 (1) You can code up what are called “sticky obstacles”, i.e., if the robot enters  
 16 an obstacle, then it stays there forever while incurring the obstacle cost at  
 17 each time instant.
- 18 (2) It is easiest to think of the runtime cost in this problem as a function of three  
 19 quantities  $q(x, u, x')$  where  $x$  is the current state,  $u$  is the control and  $x'$  is  
 20 the next state. The Bellman equation then becomes

$$J^*(x) = \min_{u \in U} \mathbb{E}_{x'} [q(x, u, x') + \gamma J^*(x')] .$$

21 You will submit your own code for this problem; there is no auto-grader.

22 **Problem 2 (Simultaneous Localization and Mapping (SLAM) with a particle**  
 23 **filter, 60 points.)** In this problem, we will implement mapping and localization in  
 24 an indoor environment using information from an IMU and a LiDAR sensor. We  
 25 have provided you data collected from a humanoid named THOR that was built at  
 26 Penn and UCLA  
 27 (<https://www.youtube.com/watch?v=JhWYYuba1nE>). You can read more about the  
 28 hardware in this paper (<https://ieeexplore.ieee.org/document/7057369>.)



1

2     **Hardware setup of Thor** The humanoid has a Hokuyo LiDAR sensor ([https://hokuyo-](https://hokuyo-usa.com/products/lidar-obstacle-detection)  
3     [usa.com/products/lidar-obstacle-detection](https://hokuyo-usa.com/products/lidar-obstacle-detection) on its head (the final version of the robot  
4     had it in its chest but this is a different version); details of this are in the code (which  
5     will be explained shortly). This LiDAR is a planar LiDAR sensor and returns 1080  
6     readings at each instant, each reading being the distance of some physical object  
7     along a ray that shoots off at an angle between  $(-135, 135)$  degrees with discretization  
8     of 0.25 degrees in an horizontal plane (shown as white rays in the picture). We will  
9     use the position and orientation of the head of the robot to calculate the orientation  
10    of the LiDAR in the body frame.

11    The second kind of observations we will use pertain to the location of the robot.  
12    However, in contrast to the previous homework where we used the raw accelerometer  
13    and gyroscope readings to get the orientation, we will directly use the  $(x, y, \theta)$  pose  
14    of the robot in the world coordinates ( $\theta$  denotes yaw). These poses were created  
15    presumably on the robot by running a filter on the IMU data (such estimates are  
16    called odometry estimates), and just as you saw some tracking errors in the previous  
17    homework, these poses will not be extremely accurate. However, we will treat them  
18    conceptually the same way as we treated Vicon in the previous homework, namely  
19    as a much more precise estimate of the pose of the robot that is used to check how  
20    well SLAM is working.

21    **Coordinate frames** The body frame is at the top of the head (X axis pointing  
22    forwards, Y axis pointing left and Z axis pointing upwards), the top of the head is at  
23    a height of 1.263m from the ground. The transformation from the body frame to the  
24    LiDAR frame depends upon the angle of the head (pitch) and the angle of the neck  
25    (yaw) and the height of the LiDAR above the head (which is 0.15m). The world  
26    coordinate frame where we want to build the map has its origin on the ground plane,

1 i.e., the origin of the body frame is at a height of 1.263m with respect to the world  
2 frame at location  $(x, y, \theta)$ .

### 3 **Data and code**

4 (a) **(0 points)** We have provided you 4 datasets corresponding to 4 different  
5 trajectories of the robot in Towne Building at Penn. For example, dataset 0 consists  
6 of two files `data/train/train_lidar0.mat` and `data/train/train_joint0.mat` which contain  
7 the LiDAR readings and joint angles respectively. The functions `load_lidar_data`  
8 and `load_joint_data` inside `load_data.py` read the data. You can run the function  
9 `show_lidar` to see the LiDAR data. Each of the data reading functions returns a  
10 data-structure where  $t$  refers to the time-stamp (in seconds) of the data, `xyth` refers  
11 to  $(x, y, \theta)$  *pose of the LiDAR* and `rpy` refers to Euler angles (roll, pitch, yaw). The  
12 joint data contains a number of fields, but we are only interested in the angle of the  
13 head and the neck at a particular time-stamp. The array `slam_t.joint.head_angles`  
14 contains the angles of neck and head respectively in the first two rows. This data is  
15 the same as the data inside the first two rows `slam_t.joint.pos` (that array contains  
16 the angles of all joints). You should read these functions carefully and check  
17 the values returned by them. The dicts `joint_names` and `joint_names_to_index`  
18 can be used to read off the data of a specific joint (we only need the head and the neck).

19  
20 (b) **(0 points)** Next look at the `slam.py` file provided to you. Read the code  
21 for the class `map_t` and `slam_t` and the comments provided in the code very  
22 carefully. You are in charge of filling in the missing pieces marked as `TODO:`  
23 `XXXXXX`. A suggested order for studying this code is as follows: `slam_t.read_data`,  
24 `slam_t.init_sensor_model`, `slam_t.init_particles`, `slam_t.rays2world`, `map_t.__init__`,  
25 `map_t.grid_cell_from_xy`. Next, the file `utils.py` contains a few standard rigid-body  
26 transformations that you will need. You should pay attention to the functions  
27 `smart_plus_2d` and `smart_minus_2d` that will be used to code up the dynamics  
28 propagation step of the particle filter.

29  
30 (c) **(10 points, dynamics step)** Next look at `main.py` which has two functions  
31 `run_dynamics_step` and `run_observation_step` which act as test functions to check  
32 if the particle filter and occupancy grid update has been updated correctly. The  
33 `run_dynamics` function plots the trajectory of the robot (as given by its IMU data in  
34 the LiDAR data-structure). It also initializes 3 particles and plots all particles at  
35 different time-steps while performing the dynamics step with a very small dynamics  
36 noise; this is a very neat way of checking if dynamics propagation in the particle  
37 filter is working correctly. This function will create two plots, one for the odometry  
38 trajectory and one more for the particle trajectories, both these trajectories should  
39 match after you code up the dynamics function `slam_t.dynamics_step` correctly.

40

1 (d) **(20 points, observation step)** The function `run_observation_step` is used to  
 2 perform the observation step of the particle filter to get an estimate of the location  
 3 of the robot and updates to the occupancy grid using observations from the LiDAR.  
 4 First read the comments for the function `slam_t.observation_step` carefully.  
 5 We first discuss the particle filtering part.

6 (i) Compute the head and neck position for the time  $t$ . For each particle,  
 7 assuming that that particle is indeed the true position of the robot, project  
 8 the LiDAR scan `slam_t.lidar[t]['scan']` into the world coordinates using the  
 9 `slam_t.ray2world` function. The end points of each ray tell us which cells in  
 10 the map are occupied, for each particle.

11 (ii) In order to compute the updated weights of the particle, we need to know  
 12 the likelihood of LiDAR scans given the state (our current occupancy grid  
 13 in the case of SLAM). We are going to use a simple model to do so

$$\log P(\text{LiDAR scan as if the robot is at particle } p \mid m) = \sum_{ij \in O} m_{ij} \quad (1)$$

14 where  $O$  is the set of occupied cells as detected by the LiDAR scan assuming  
 15 the robot is at particle  $p$  and  $m_{ij}$  is our current estimate of the binarized  
 16 map (more on this below). In simple words, if the occupied cells as given  
 17 by our LiDAR match the occupied cells in the binarized map created from  
 18 the past observations, then we say the log-probability of particle  $p$  is large.

19 (iii) You will next implement the function `slam_t.update_weights` that takes  
 20 the log-probability of each particle  $p$ , its previous weights, calculates the  
 21 updated weights of the particles.

22 (iv) Typically, resampling step (`slam_t.stratified_resampling`) is performed only if  
 23 the effective number of particles (as computed in `slam_t.resample_particles`)  
 24 falls below a certain threshold (30% in the code). Implement resampling as  
 25 we discussed in the lecture notes.

26 **Mapping** We have a number of particles  $p^i = (x^i, y^i, \theta^i)$  that together give  
 27 an estimate of the distribution of the location of the robot. For this homework,  
 28 you will only use the particle with the largest weight to update the map although  
 29 typically we update the map using all particles. Our goal is simple: we want to  
 30 increase `map_t.log_odds` array at cells that are recorded as obstacles by the LiDAR  
 31 and decrease the values in all other cells. You should add `slam_t.log_odds_occ`  
 32 to all occupied cells and add `slam_t.log_odds_free` from all cells in the map. It is  
 33 also a good idea to clip the `log_odds` to like between `[-slam_t.map.log_odds_max,`  
 34 `slam_t.map.log_odds_max]` to prevent increasingly large values in the `log_odds`  
 35 array. The array `slam_t.map.cells` is a binarized version of the map (which is used  
 36 above to calculate the observation likelihood).


1 Check the `run_observation_step` function after you have implemented the obser-  
2 vation step.

3 (e) Since the map is initialized to zero at the beginning of SLAM which results in  
4 all observation log-likelihoods to be zero in (1), we need to do something special  
5 for the first step. We will use the first entry in `slam_t.lidar[0]['xyth']` to get an  
6 accurate pose for the robot and use its corresponding LiDAR readings to initialize  
7 the occupancy grid. You can do this easily by initializing the particle filter to have  
8 just one particle and simply calling the `slam_t.observation_step` as shown in `main.py`.

9 (f) **(30 points)** You will now run the full SLAM algorithm that performs one  
10 dynamics step and observation step at each iteration in the function `run_slam` in  
11 `main.py`. Make sure to start SLAM only after the time when you have both LiDAR  
12 scans and joint readings (the two arrays start at different times). For all 4 datasets,  
13 you will plot the final binarized version of the map,  $(x, y)$  location of the particle  
14 in the particle filter with the largest weight at each time-step and the odometry  
15 trajectory  $(x, y)$  (in a different color); this counts for 10 points each.

16 **Some Notes** This problem is much easier and shorter than it may seem. You  
17 should go through these steps carefully and in the suggested order. You should make  
18 sure that the results of the previous step are correct before proceeding. The two  
19 functions in `main.py` to check the dynamics and observation step are very important  
20 to find bugs. You do not need to implement more than 100 lines of code.

21 **Problem 3 (Swinging up an Acrobot, 30 points).** The Acrobot is a classical  
22 under-actuated dynamical system. It is a model for a gymnast swinging up on a bar.




23

24 The gymnast only applies control input at his waist. He begins from his initial  
25 position of hanging down from the bar and wants to end up at a terminal position  
26 where he is upright. Note that the vertical position is an equilibrium position,  
27 and therefore, the gymnast will stay there indefinitely without any control input  
28 (<https://www.youtube.com/watch?v=O2b03YtMeRU>). In this problem, you will  
29 write an LQR controller to take the Acrobot from the initial condition to the terminal  
30 condition. It is difficult to directly use LQR for such a highly nonlinear system. We  
31 will therefore use something called as an energy shaping controller to first swing



- 1 up the Acrobot to the near vertical position and then switch on the LQR controller  
 2 which will keep the robot in the upright position (which is the equilibrium point).



3

- 4 The equations of motion for an Acrobot are as follows. You can read  
 5 <http://underactuated.mit.edu/acrobot.html> to see how they are derived by writing  
 6 down the Lagrangian. Let  $q(t) = [q_1(t), q_2(t)]$  be the joint angles.

$$M(q)\ddot{q} + C(q, \dot{q})\dot{q} + G(q) = Bu \quad (2)$$

7 where

$$M(q) = \begin{bmatrix} I_1 + I_2 + m_2 l_1^2 + m_2 l_1 l_2 \cos q_2 & \frac{1}{2} I_2 + m_2 l_1 l_2 \cos q_2 \\ I_2 + \frac{1}{2} m_2 l_1 l_2 \cos q_2 & I_2 \end{bmatrix}$$

$$C(q, \dot{q}) = \begin{bmatrix} -m_2 l_1 l_2 \sin q_2 \dot{q}_2 & -\frac{1}{2} m_2 l_1 l_2 \sin q_2 \dot{q}_2 \\ \frac{1}{2} m_2 l_1 l_2 \sin q_2 \dot{q}_1 & 0 \end{bmatrix}$$

$$G(q) = \begin{bmatrix} (\frac{1}{2} m_1 l_1 g + m_2 l_1 g) \sin q_1 + \frac{1}{2} m_2 l_2 g \sin (q_1 + q_2) \\ \frac{1}{2} m_2 l_2 g \sin (q_1 + q_2) \end{bmatrix}$$

$$B = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

- 8 **(a) Linearize the dynamics and implement LQR (10 points)** For the second  
 9 order dynamics, define the state  $x \equiv [q, \dot{q}] \in \mathbb{R}^4$ . We will first linearize (2) around  
 10 the vertical position  $x_f = [\pi, 0, 0, 0]^\top$ . This involves defining a new state

$$\delta x = x - x_f$$

- 11 and writing down (2) in terms of the new state  $\delta x$ .

$$\delta \dot{x} = A \delta x + B u.$$

- 12 Write down your  $A$  and  $B$  matrices and how you calculated them. You can use  
 13 Mathematica or Matlab or Sympy (show the code if you use these softwares) to take  
 14 derivatives if you'd like but doing it by hand is quicker and easier; just make sure  
 15 to double check your calculations. You can verify your linearized dynamics using

1 finite differences. That's:

$$\frac{\partial f}{\partial x_i}(x_f) \approx \frac{f(x_f + \epsilon \mathbf{1}_i) - f(x_f)}{\epsilon}$$

2 for some small  $\epsilon$ .

3 You will now choose matrices  $Q$  and  $R$  to set the cost function to be

$$\frac{1}{2} \int_0^\infty dt \delta x(t)^\top Q \delta x(t) + u(t)^\top R u(t).$$

4 We know that for infinite-horizon LQR, the optimal cost-to-go is quadratic in the  
5 state, i.e.,  $V(\delta x) = \delta x^\top P \delta x$ . The matrix  $P$  can be obtained by solving a Riccati  
6 equation:

$$A^\top P + PA - PBR^{-1}B^\top P + Q = 0$$

7 and the corresponding optimal feedback control is given by


$$u = -K\delta x$$

8 where the Kalman gain is  $K = R^{-1}B^\top P$ .

9 In the code that we have provided, you will fill in the linearized dynamics  
10 in `get_lqr_term` and the LQR controller in `get_control_efforts`. You should use  
11 `scipy.linalg.solve_continuous_are` to solve the Riccati equation.

12 To check your controller, you should try a few initial conditions close to the  
13 vertical position, e.g.,  $x_0 = [\pi - 0.1, 0, 0.2, 0]$ . Your controller should be able to  
14 balance the Acrobot to the equilibrium point.

15 **(b) Check initial conditions for which LQR works (5 points)** If you try a few  
16 others, you will notice that the LQR controller is not able to stabilize the robot.  
17 Argue why or why not. There is however a neat trick that we can use to calculate  
18 the set of initial conditions that LQR works for. Notice that the value function is  
19  $V(\delta x) = \delta x^\top P \delta x$ . For write down the  $P$  matrix that you have calculated in the  
20 previous step and show the eigenvalues of this matrix. You will notice that the  
21 eigenvalues are not all equal, some are very large while some are rather small. This  
22 indicates that the level set, i.e., the set of states with small cost to go according to  
23 LQR looks like an ellipse with a large aspect ratio:



24

25 If we choose some small value of  $c$ , then we can think of the states in the set

$$\Omega_c = \left\{ x : \delta x^\top P \delta x \leq c \right\}$$

1 as the set of states that LQR works for. Use a small value of  $c$  (there is neat way to  
 2 choose this if you think using the eigenvalues of  $P$ ) and draw a plot of  $q_1(t)$  vs  $q_2(t)$   
 3 for at least two trajectories of the Acrobot, one that starts inside this set and one that  
 4 starts outside this set.

5 **(c) Energy shaping controller (10 points)** LQR can balance the robot once it is  
 6 close to the terminal condition. We will now use a different controller to swing up  
 7 the robot from its starting state  $x_0 = [0, 0, 0, 0]$ , which is at the bottom.

8 If we have a fully-actuated system like  $\ddot{q} = u$ , it would be easy to stabilize it with  
 9 feedback control, e.g.,  $\ddot{q} = u = k_p q + k_d \dot{q}$  with negative  $k_p, k_d < 0$ . The Acrobot  
 10 is not a fully-actuated system: there is only one control input which is the torque at  
 11 the waist and there are two degrees of freedom. But we can rewrite the system in a  
 12 special way using a technique called partial feedback linearization. Here is how it  
 13 works.

$$\begin{aligned} M_{11}\ddot{q}_1 + M_{12}\ddot{q}_2 &= \tau_1 \\ M_{21}\ddot{q}_1 + M_{22}\ddot{q}_2 &= \tau_2 + u \end{aligned} \quad (3)$$

14 where  $\tau = [\tau_1, \tau_2]^T \equiv -C\dot{q} - G$  from (2). Observe that this is a linear system in  
 15 the two variables  $[\ddot{q}_1, \ddot{q}_2]$  and we can solve this system to get

$$\begin{aligned} \ddot{q}_1 &= \text{func}_1(q, \dot{q}, u) \\ \ddot{q}_2 &= \text{func}_2(q, \dot{q}, u). \end{aligned}$$

16 In our problem, the torque directly moves the second angle  $q_2$  (the first angle  $q_1$  is  
 17 actuated only indirectly by the torque due to conservation of angular momentum).  
 18 We therefore will use the second equation to obtain a relationship

$$u = \text{func}_3(\ddot{q}_2, q, \dot{q}). \quad (4)$$

19 Write down what your  $\text{func}_3$  is and explain in detail how you found it.

20 This is pretty neat because (4) tells us what control  $u$  we should use if we want  
 21 a particular  $q, \dot{q}$  and  $\ddot{q}_2$ . We will now use this calculation to calculate a particular  
 22 feedback control, the one that consistently adds energy to the robot and makes it  
 23 swing up.

24 The Acrobot has the smallest energy (sum of potential energy and kinetic energy  
 25 due to the angular velocity) when it is at the bottom. When it is stationary at the  
 26 top, it only has potential energy which is larger than the energy at the bottom. If the  
 27 Acrobot is to swing up, we need to add at least the difference between the two into  
 28 the system. The basic idea of an energy shaping controller is to realize that just like  
 29 you push down with your legs when you swing up on a swing in the park, i.e., you  
 30 push towards the direction of motion instead of against it, we can calculate a control  
 31 input that adds energy by applying a torque in the direction of motion.

$$\begin{aligned}
\tilde{E} &= E(x) - E(x_f) \\
\bar{u} &= \dot{q}_1 \tilde{E} \\
\ddot{q}_2^d &= -k_1 q_2 - k_2 \dot{q}_2 + k_3 \bar{u} \\
u &= \text{func}_3(\ddot{q}_2^d, q, \dot{q})
\end{aligned} \tag{5}$$

1 In the above equation we have chosen a desired angular acceleration  $\ddot{q}_2^d$  to calculate  
2 the actual control  $u$ , but this desired  $\ddot{q}_2^d$  is calculated in such a way that as the gap  
3 between  $E(x)$  and  $E(x_f)$  goes to zero, the second angle  $q_2$  forms a stable system  
4 by itself (for  $k_1, k_3 > 0$ ), i.e.,  $q_2 \rightarrow 0$  if  $\bar{u} = 0$  from any initial condition.

5 The most important point of the above equations is that  $\bar{u} = \dot{q}_1 \tilde{E}$ ; this is the  
6 energy shaping part. In simple words,  $\bar{u}$  is proportional to how much energy should  
7 be added/subtracted (depending upon  $\tilde{E}$ ) at each time instant. This “effective  
8 control” leads to an angular acceleration  $\ddot{q}_2^d$ , which gives us the actual control  $u$   
9 that we should take using the final equation. In the third equation, we have also  
10 added feedback in the form of  $-k_1 q_2 - k_2 \dot{q}_2$  which makes sure that if  $q_2$  is far from  
11 zero, we choose a slightly larger  $\ddot{q}_2^d$  than what would be dictated by simply the  
12 effective control term  $k_3 \bar{u}$ . You should choose the values of  $k_1, k_2, k_3 > 0$  yourself.  
13 Complete the *get\_swingup\_input* function in the code that we have provided. **When**  
14 **you implement all these formulae, do not forget to convert  $q$  back to  $[-\pi, \pi]$**   
15 **using `numpy.unwrap` every time you use it. This is very important.**

16 **(d) Run the entire system (5 points)** Plot the joint angles and the control input  
17 for when the Acrobot starts from the stationary state at the bottom, uses energy  
18 shaping to reach a state that lies within the ellipse of initial conditions that are  
19 suitable for LQR and then uses LQR to stabilize in the upright position. You find the  
20 code provided useful for animating the motion of the Acrobot for both debugging  
21 and final plotting.